

Final Question

Huafeng Zhang (217107707)

2019.7.26

1. load data, remove na, log-transform

```
library(ISLR)
hit <- Hitters #322 obs
hit <- na.omit(hit) #263 obs
hit$log.Salary <- log(hit$Salary)
#names(hit)
```

2. split to train/test

```
tr.hit <- hit[1:200,]
te.hit <- hit[201:263,]
```

3. Random Forest

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.6.1
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(345)
```

```
#25 trees
```

```
rf25 <- randomForest(log.Salary ~ . - Salary, data = tr.hit, ntree = 25, Importance = T)
```

```
ps25 <- predict(rf25, newdata = te.hit)
```

```
mse25 <- mean((te.hit$log.Salary-ps25)^2)
```

```
#100 trees
```

```
#rf100 <- grow(rf25, 75)
```

```
rf100 <- randomForest(log.Salary ~ . - Salary, data = tr.hit, ntree = 100, Importance = T)
```

```
ps100 <- predict(rf100, newdata = te.hit)
```

```
mse100 <- mean((te.hit$log.Salary-ps100)^2)
```

```
#500 trees
```

```
#rf500 <- grow(rf100, 400)
```

```
rf500 <- randomForest(log.Salary ~ . - Salary, data = tr.hit, ntree = 500, Importance = T)
```

```
ps500 <- predict(rf500, newdata = te.hit)
```

```

mse500 <- mean((te.hit$log.Salary-ps500)^2)

#1000 trees
#rf1000 <- grow(rf500, 500)
rf1000 <- randomForest(log.Salary ~ . - Salary, data = tr.hit, ntree = 1000, Importance = T)
ps1000 <- predict(rf1000, newdata = te.hit)
mse1000 <- mean((te.hit$log.Salary-ps1000)^2)

print(paste('MSE of 25 trees is', mse25))

## [1] "MSE of 25 trees is 0.227011143697994"

print(paste('MSE of 100 trees is', mse100))

## [1] "MSE of 100 trees is 0.228650313687036"

print(paste('MSE of 500 trees is', mse500))

## [1] "MSE of 500 trees is 0.216076387586862"

print(paste('MSE of 1000 trees is', mse1000))

## [1] "MSE of 1000 trees is 0.217779438384326"

###Given the 4 results, using 500 trees gives the best result with the smallest MSE of 0.216.

```

4. identify the most important variables

```

print(rf500) #No. of variables tried at each split: 6

##
## Call:
## randomForest(formula = log.Salary ~ . - Salary, data = tr.hit,      ntree = 500, Importance = T)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 6
##
##              Mean of squared residuals: 0.2135731
##              % Var explained: 74.33

importance(rf500, type=2)

##              IncNodePurity
## AtBat              6.0610657
## Hits              5.8677837

```

```
## HmRun      1.9309252
## Runs       4.5674917
## RBI        4.8661721
## Walks      5.0550154
## Years      7.2652886
## CAtBat     34.7331263
## CHits      24.5292756
## CHmRun     5.5484426
## CRuns      26.1100611
## CRBI       15.1346342
## CWalks     16.3724500
## League     0.1046373
## Division   0.1991141
## PutOuts    2.7902861
## Assists    1.4517977
## Errors     1.2341805
## NewLeague  0.2172757
```

The most important variables (with the largest values) associated with predicting Salary are 'CAtBat', 'CHits', 'CRuns', 'CWalks', and 'CRBI'.

5. Boosting and plot of learning rates vs. train MSEs

```
library(gbm)
```

```
## Warning: package 'gbm' was built under R version 3.6.1
```

```
## Loaded gbm 2.1.5
```

```
set.seed(567)
```

```
#0.01 learning rate
```

```
b01 <- gbm(formula = log.Salary ~ . - Salary, data = tr.hit,
           distribution = "gaussian", n.trees = 1000,
           interaction.depth = 1, shrinkage = 0.01, verbose = F)
trps01 <- predict(b01, n.trees = b01$n.trees, newdata = tr.hit)
trmse01 <- mean((tr.hit$log.Salary-trps01)^2)
```

```
#0.001 learning rate
```

```
b001 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
            distribution = "gaussian", n.trees = 1000,
            interaction.depth = 1, shrinkage = 0.001, verbose = F)
trps001 <- predict(b001, n.trees = b001$n.trees, newdata = tr.hit)
trmse001 <- mean((tr.hit$log.Salary-trps001)^2)
```

```
#0.1 learning rate
```

```
b1 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
          distribution = "gaussian", n.trees = 1000,
          interaction.depth = 1, shrinkage = 0.1, verbose = F)
trps1 <- predict(b1, n.trees = b1$n.trees, newdata = tr.hit)
```

```

trmse1 <- mean((tr.hit$log.Salary-trps1)^2)

#0.2 learning rate
b2 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
          distribution = "gaussian", n.trees = 1000,
          interaction.depth = 1, shrinkage = 0.2, verbose = F)
trps2 <- predict(b2, n.trees = b2$n.trees, newdata = tr.hit)
trmse2 <- mean((tr.hit$log.Salary-trps2)^2)

print(paste('train MSE of 0.01 learning rate at d=1 is', trmse01))

## [1] "train MSE of 0.01 learning rate at d=1 is 0.138358148208026"

print(paste('train MSE of 0.001 learning rate at d=1 is', trmse001))

## [1] "train MSE of 0.001 learning rate at d=1 is 0.361502516103347"

print(paste('train MSE of 0.1 learning rate at d=1 is', trmse1))

## [1] "train MSE of 0.1 learning rate at d=1 is 0.0486036557696533"


print(paste('train MSE of 0.2 learning rate at d=1 is', trmse2))

## [1] "train MSE of 0.2 learning rate at d=1 is 0.0239267402963262"

#plot (train)
plot(c(0.001,0.01,0.1,0.2),
     c(trmse001, trmse01, trmse1, trmse2),
     type = 'l', col = 'coral3',
     main = 'Learning Rate vs. Train MSE at d=1',
     xlab = 'Learning Rate',
     ylab = 'Train MSE')

```

Learning Rate vs. Train MSE at d=1



6. plot of learning rates vs. test MSEs

```
#interaction.depth = 1
library(gbm)
set.seed(567)

#0.01 learning rate
ps01 <- predict(b01, n.trees = b01$n.trees, newdata = te.hit)
mse01 <- mean((te.hit$log.Salary-ps01)^2)

#0.001 learning rate
ps001 <- predict(b001, n.trees = b001$n.trees, newdata = te.hit)
mse001 <- mean((te.hit$log.Salary-ps001)^2)

#0.1 learning rate
ps1 <- predict(b1, n.trees = b1$n.trees, newdata = te.hit)
mse1 <- mean((te.hit$log.Salary-ps1)^2)

#0.2 learning rate
ps2 <- predict(b2, n.trees = b2$n.trees, newdata = te.hit)
mse2 <- mean((te.hit$log.Salary-ps2)^2)

print(paste('MSE of 0.01 learning rate at d=1 is', mse01))
```

```
## [1] "MSE of 0.01 learning rate at d=1 is 0.281627516952774"
```

```
print(paste('MSE of 0.001 learning rate at d=1 is', mse001))
```

```
## [1] "MSE of 0.001 learning rate at d=1 is 0.337744738172609"
```


```
print(paste('MSE of 0.1 learning rate at d=1 is', mse1))
```

```
## [1] "MSE of 0.1 learning rate at d=1 is 0.253919209796006"
```

```
print(paste('MSE of 0.2 learning rate at d=1 is', mse2))
```

```
## [1] "MSE of 0.2 learning rate at d=1 is 0.279011819536214"
```

```
#plot (test)  
plot(c(0.001,0.01,0.1,0.2),  
      c(mse001,mse01,mse1,mse2),  
      type = 'l', col = 'coral3',  
      main = 'Learning Rate vs. Test MSE at d=1',  
      xlab = 'Learning Rate',  
      ylab = 'Test MSE')
```



7. try different interaction depths

```
#interaction.depth = 2
set.seed(666)

#0.01 learning rate
b01 <- gbm(formula = log.Salary ~ . - Salary, data = tr.hit,
            distribution = "gaussian", n.trees = 1000,
            interaction.depth = 2, shrinkage = 0.01, verbose = F)
ps01 <- predict(b01, n.trees = b01$n.trees, newdata = te.hit)
mse01 <- mean((te.hit$log.Salary-ps01)^2)

#0.001 learning rate
b001 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
            distribution = "gaussian", n.trees = 1000,
            interaction.depth = 2, shrinkage = 0.001, verbose = F)
ps001 <- predict(b001, n.trees = b001$n.trees, newdata = te.hit)
mse001 <- mean((te.hit$log.Salary-ps001)^2)

#0.1 learning rate
b1 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
           distribution = "gaussian", n.trees = 1000,
           interaction.depth = 2, shrinkage = 0.1, verbose = F)
ps1 <- predict(b1, n.trees = b1$n.trees, newdata = te.hit)
mse1 <- mean((te.hit$log.Salary-ps1)^2)

#0.2 learning rate
b2 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
           distribution = "gaussian", n.trees = 1000,
           interaction.depth = 2, shrinkage = 0.2, verbose = F)
ps2 <- predict(b2, n.trees = b2$n.trees, newdata = te.hit)
mse2 <- mean((te.hit$log.Salary-ps2)^2)

print(paste('MSE of 0.01 learning rate at d=2 is', mse01))
```

```
## [1] "MSE of 0.01 learning rate at d=2 is 0.276960632061591"
```

```
print(paste('MSE of 0.001 learning rate at d=2 is', mse001))
```

```
## [1] "MSE of 0.001 learning rate at d=2 is 0.314604324482154"
```

```
print(paste('MSE of 0.1 learning rate at d=2 is', mse1))
```


```
## [1] "MSE of 0.1 learning rate at d=2 is 0.304501837582825"
```

```
print(paste('MSE of 0.2 learning rate at d=2 is', mse2))
```

```
## [1] "MSE of 0.2 learning rate at d=2 is 0.308179134663943"
```

```
#plot (test)
plot(c(0.001,0.01,0.1,0.2),
     c(mse001,mse01,mse1,mse2),
     type = 'l', col = 'coral3',
     main = 'Learning Rate vs. Test MSE at d=2',
     xlab = 'Learning Rate',
     ylab = 'Test MSE')
```

Learning Rate vs. Test MSE at d=2



```
#interaction.depth = 4
set.seed(233)

#0.01 learning rate
b01 <- gbm(formula = log.Salary ~ . - Salary, data = tr.hit,
           distribution = "gaussian", n.trees = 1000,
           interaction.depth = 4, shrinkage = 0.01, verbose = F)
ps01 <- predict(b01, n.trees = b01$n.trees, newdata = te.hit)
mse01 <- mean((te.hit$log.Salary-ps01)^2)

#0.001 learning rate
b001 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
            distribution = "gaussian", n.trees = 1000,
            interaction.depth = 4, shrinkage = 0.001, verbose = F)
ps001 <- predict(b001, n.trees = b001$n.trees, newdata = te.hit)
mse001 <- mean((te.hit$log.Salary-ps001)^2)
```



```
#0.1 learning rate
b1 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
          distribution = "gaussian", n.trees = 1000,
          interaction.depth = 4, shrinkage = 0.1, verbose = F)
ps1 <- predict(b1, n.trees = b1$n.trees, newdata = te.hit)
mse1 <- mean((te.hit$log.Salary-ps1)^2)
```

```
#0.2 learning rate
b2 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
          distribution = "gaussian", n.trees = 1000,
          interaction.depth = 4, shrinkage = 0.2, verbose = F)
ps2 <- predict(b2, n.trees = b2$n.trees, newdata = te.hit)
mse2 <- mean((te.hit$log.Salary-ps2)^2)
```

```
print(paste('MSE of 0.01 learning rate at d=4 is', mse01))
```

```
## [1] "MSE of 0.01 learning rate at d=4 is 0.280471018407855"
```

```
print(paste('MSE of 0.001 learning rate at d=4 is', mse001))
```

```
## [1] "MSE of 0.001 learning rate at d=4 is 0.307862913776067"
```

```
print(paste('MSE of 0.1 learning rate at d=4 is', mse1))
```


```
## [1] "MSE of 0.1 learning rate at d=4 is 0.288221061450338"
```

```
print(paste('MSE of 0.2 learning rate at d=4 is', mse2))
```

```
## [1] "MSE of 0.2 learning rate at d=4 is 0.263558950473661"
```

```
#plot (test)
plot(c(0.001,0.01,0.1,0.2),
     c(mse001,mse01,mse1,mse2),
     type = 'l', col = 'coral3',
     main = 'Learning Rate vs. Test MSE at d=4',
     xlab = 'Learning Rate',
     ylab = 'Test MSE')
```


Learning Rate vs. Test MSE at d=4



For interaction depth of 1, $\lambda = 0.1$ gives the smallest test MSE of 0.253919209796006.
For interaction depth of 2, $\lambda = 0.01$ gives the smallest test MSE of 0.276960632061591.
For interaction depth of 4, $\lambda = 0.2$ gives the smallest test MSE of 0.263558950473661.

8. most important variables from boosted tree

```
set.seed(567) #the same seed as above for d=1
b1 <- gbm(log.Salary ~ . - Salary, data = tr.hit,
          distribution = "gaussian", n.trees = 1000,
          interaction.depth = 1, shrinkage = 0.1, verbose = F)
summary(b1)
```



##	var	rel.inf
## CWalks	CWalks	15.6449722
## CATBat	CATBat	12.9375051
## CRBI	CRBI	10.1040118
## PutOuts	PutOuts	8.0896231
## CRuns	CRuns	6.5267876
## Walks	Walks	6.1674450
## CHmRun	CHmRun	5.5898294
## Years	Years	5.3205366
## Hits	Hits	4.9848118
## RBI	RBI	4.8630142
## Assists	Assists	4.6119239
## HmRun	HmRun	3.5604525
## AtBat	AtBat	3.0374210
## Runs	Runs	2.8937549
## Errors	Errors	2.3637443
## CHits	CHits	2.1652687
## Division	Division	0.5091896
## NewLeague	NewLeague	0.3701619
## League	League	0.2595463

Given the result, 'CWalks', 'CAAtBat', and 'CRBI' are the most important variables.

9. compare test MSEs from RF and boosting

In random forest, using 500 trees gives the best result with the smallest MSE of 0.216, which is better than any of the test MSE from boosting.